



DEPARTMENT OF INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

**Implementation of Attribute-Based
Encryption in Rust on ARM Cortex M
Processors**

Daniel Bücheler



DEPARTMENT OF INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

**Implementation of Attribute-Based
Encryption in Rust on ARM Cortex M
Processors**

**Implementierung von Attributbasierter
Verschlüsselung in Rust auf ARM Cortex
M Prozessoren**

Author:	Daniel Bücheler
Supervisor:	Prof. Dr. Claudia Eckert
Advisor:	Stefan Hristozov
Submission Date:	15.04.2021

I confirm that this bachelor's thesis in informatics is my own work and I have documented all sources and material used.

Munich, 15.04.2021

Daniel Bücheler

Acknowledgments

Abstract

Contents

Acknowledgments	iii
Abstract	iv
1 Introduction	1
1.1 Section	1
1.1.1 Subsection	1
2 Background	3
2.1 Classic Symmetric and Asymmetric Cryptography	3
2.2 Attribute-Based Encryption	4
2.2.1 KP-ABE and CP-ABE	4
2.2.2 Components of an ABE scheme	5
2.3 Shamir's Secret Sharing	5
2.4 Elliptic Curves	7
2.4.1 Point Addition	7
2.4.2 Groups on Elliptic Curves	8
2.4.3 Hardness Assumptions	9
2.4.4 Bilinear Pairings	9
2.5 Different ABE schemes for use in Embedded Devices	10
3 Challenges encountered during the implementation	11
3.1 Rust specialties	11
3.2 Lack of Operating System	11
3.3 Performance Limitations	12
List of Figures	13
List of Tables	14
Bibliography	15

1 Introduction

1.1 Section

Citation test [**latex**].

1.1.1 Subsection

See Table 1.1, Figure 1.1, Figure 1.2, Figure 1.3.

Table 1.1: An example for a simple table.

A	B	C	D
1	2	1	2
2	3	2	3



Figure 1.1: An example for a simple drawing.

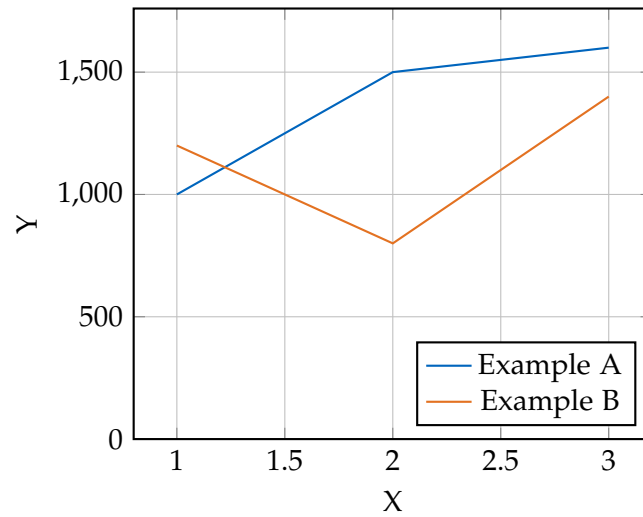


Figure 1.2: An example for a simple plot.

```
SELECT * FROM tbl WHERE tbl.str = "str"
```

Figure 1.3: An example for a source code listing.

2 Background

This chapter shall provide a high-level introduction to the cryptographical and mathematical tools used to implement attribute-based encryption in this thesis. For further reference, please refer to

2.1 Classic Symmetric and Asymmetric Cryptography

Today's conventional cryptography knows two main classes of cryptosystems: *Symmetric* or *Private-Key* systems and *Asymmetric* or *Public-Key* systems. The main difference lies in their use of encryption and decryption keys:

In *symmetric* systems, the key used for encryption and decryption is identical. That is, a user *Alice* encrypting a message to send to another user *Bob* will encrypt the message using a key k that she had agreed on with Bob beforehand. When he receives the encrypted message, Bob will use the same key k to decrypt it. See Figure 2.1a.

In *asymmetric* systems, on the other hand, the keys used for encryption and decryption differ. When Alice encrypts a message with key k_{enc} , she will not be able to decrypt it again. Instead, when Bob receives the encrypted message, he will use a different key k_{dec} to decrypt it. See Figure 2.1b. Thus, in symmetric systems, keys always come in *pairs* of an *encryption key* k_{enc} and a *decryption key* k_{dec} . Because k_{enc} can not be used to decrypt messages meant for Bob, there is no harm to making it publicly available. For example, he might put it up on his website for anyone wishing to send him an encrypted message to download. This is why k_{enc} is also often called the *public key* and k_{dec} the *private key*.

Asymmetric cryptosystems make secure communication among a large group of participants much easier: Consider n participants wanting to communicate securely using a symmetric system. Each participant would need to share a unique secret key with each of the other participants, requiring a total of $\frac{n(n-1)}{2}$ keys. In the asymmetric setting, one key per participant is sufficient: The same public key may be shared with the whole group, as the private key remains private anyway. This reduces the total number of keys to n .

Another problem remains, however: Encrypting a single message to a large number of participants requires encrypting it with everyone's public key separately. For a large




Figure 2.1: Keys used for encryption and decryption in different classes of encryption schemes

number of recipients, this is inefficient. So, for example, to encrypt a message for all students of a certain university, we'd need to obtain each student's public key and encrypt the message with each key separately.

Even worse, what if we want to encrypt data for any student of said university, even if they *haven't joined the university yet*. In this case, our only option using classic public-key cryptography would be to have some trusted instance encrypt the data for any new student after they joined the university. Attribute-Based encryption solves this problem much more nicely.

2.2 Attribute-Based Encryption

Attribute-Based Encryption Schemes (ABE schemes) are asymmetric in the sense that different keys are used for encryption and decryption. However, in contrast to classic asymmetric systems, the 'public key' used for encryption is not attached to an identity, but to certain attributes defined by the system. This is represented by a tree of attributes in Figure 2.1c.

2.2.1 KP-ABE and CP-ABE

ABE schemes can further be divided into *Key-Policy* and *Ciphertext-Policy* schemes (KP-ABE and CP-ABE, respectively). In KP-ABE, the ciphertext is associated with a

set of attributes, and the key is associated with an access structure. CP-ABE works the other way around, so the key is associated with a set of attributes and the ciphertext is encrypted under an access structure.

In both cases, a ciphertext can be decrypted if and only if the set of attributes specified in one part satisfy the access policy associated with the other part.

2.2.2 Components of an ABE scheme

With respect to ABE, we say that a ciphertext is *encrypted under an access structure* (under a set of attributes, respectively) if it has been encrypted using the public key derived from the access structure (set of attributes, resp.). Because the public key w.r.t. some access structure (set of attributes, resp.) can be derived by anyone, private keys are usually referred to as just *keys*.

An Attribute-Based Encryption scheme (ABE scheme) consists of the following four algorithms:

- *Setup*. Run once by the central trusted authority (key generation center). Sets up the system and outputs a public master key PK and a private master secret s .
- *KeyGen*(PK, s, ω). Run exclusively by the trusted authority. For a given set of attributes (CP-ABE) or access structure (KP-ABE) ω , uses the master secret s to generate a private key k corresponding to the given access structure.
- *Encrypt*(PK, m, ω). Run by any participant of the system. On input of a plaintext message m and an access structure (CP-ABE) or set of attributes (KP-ABE) ω , outputs a ciphertext c .
- *Decrypt*(c, k). Run by any participant holding a private key generated by *KeyGen*. On input of a ciphertext c created by *Encrypt* and a key k generated by *KeyGen*, it outputs a correctly decrypted message if and only if
 - in the case of CP-ABE: the access structure under which m was encrypted is satisfied by the set of attributes under which k was created
 - in the case of KP-ABE: the set of attributes under which m was encrypted satisfies the access structure under which k was created.

2.3 Shamir's Secret Sharing

This secret sharing scheme based on polynomial interpolation was first introduced by Adi Shamir in 1979 [1]. It allows a secret s , which is generally just a number, to be shared among a number of n participants. The shares are computed such that s can




Figure 2.2: Example for a $(5,4)$ -threshold scheme with $s = 12$ and $p(x) = 12 + 7x - 6x^2 + x^3$. As $p(x)$ has degree three, at least four of the five shares are required to reconstruct s .

be reconstructed if, and only if, at least k participants meet and combine their shares. Such a theme is then called a (k, n) -threshold scheme. [1]

Shamir's scheme makes use of a property of polynomials: A polynomial of degree d is unambiguously determined by $d + 1$ points (x_i, y_i) . In other words, any polynomial of degree d can be unambiguously interpolated (reconstructed) from $d + 1$ distinct points. Thus, if we hide our secret in a polynomial $p(x)$ of degree d and give out secret shares as $(x, p(x))$ with $x = 1, 2, 3, \dots$, we can use $d + 1$ shares to reconstruct $p(x)$ and by that the secret.[1]

That is, to create a simple (k, n) -threshold scheme to share a secret s , do the following: [1]

1. Pick coefficients a_1, \dots, a_{k-1} at random
2. Set $a_0 = s$. This results in the polynomial $p(x) = a_0 + a_1x + \dots + a_{k-1}x^{k-1}$. Note that $p(0) = s$.
3. The secret shares are $(1, p(1)), (2, p(2)), \dots, (n, p(n))$. Give one to each participant.

To reconstruct the secret from any subset of k shares, simply interpolate the polynomial $p(x)$ and evaluate $p(0) = s$ [1]. See also Figure 2.2 for illustration.

The polynomial can be interpolated unambiguously if at least k points are known, because it has degree $k - 1$. If only $k - 1$ points (or less) are known, there are infinitely

many ways to construct a polynomial that passed through all $k - 1$ points. Therefore, then there are also infinitely many possibilities for s . [1]

2.4 Elliptic Curves

The mathematics of modern cryptosystems (including, but not limited to ABE) work on a great variety of mathematical structures, and elliptic curves are just one of them. They have become very popular since their discovery by *citation needed in year needed* because they provide an equivalent security level at shorter key length and smaller computational cost than other systems (e.g. based on RSA) [2]. Exact definitions and notations differ, these are taken from the textbook *Introduction to Modern Cryptography* by Katz and Lindell [2].

Elliptic curves (over finite fields) are defined by equations of the following form, where $p \leq 5$ prime: [2]

$$y^2 = x^3 + ax + b \pmod{p} \quad (2.1)$$

a and b are the curve parameters, and it is required that $4a^2 + 27b^2 \not\equiv 0 \pmod{p}$ (this makes sure there are no repeated roots).[2]

Then, the elliptic curve $E(\mathbb{Z}_p)$ is the set of coordinates $(x, y) \in \mathbb{Z}_p^2$ that satisfies the equation above. A special value \mathcal{O} is added to denote the *point at infinity*. [2]

$$E(\mathbb{Z}_p) := \{(x, y) | x, y \in \mathbb{Z}_p \text{ and } y^2 = x^3 + ax + b \pmod{p}\} \cup \{\mathcal{O}\} \quad (2.2)$$

2.4.1 Point Addition

Now, it is possible to show that every line intersecting a curve $E(\mathbb{Z}_p)$ intersects it in exactly three points, if you (1) count tangential intersections double and (2) count any vertical line as intersecting the curve in the point at infinity \mathcal{O} . Therefore, \mathcal{O} can be thought of as sitting 'above' the end of the y -axis.[2]

Using this intersecting line, we can define an operation on curve points called (*point*) *addition* and written using $+$: [2]

- \mathcal{O} is the identity, i.e. for any $P \in E(\mathbb{Z}_p)$, $P + \mathcal{O} = \mathcal{O} = \mathcal{O} + P$
- For two points $P_1, P_2 \neq \mathcal{O}$, their sum $P_1 + P_2$ is evaluated by drawing the line through P_1 and P_2 (if $P_1 = P_2$, draw a tangential line). This line will intersect the curve in a third point, $P_3 = (x_3, y_3)$. Then the result of the addition is $P_1 + P_2 = (x_3, -y_3)$, i.e. P_3 is reflected in the x -axis. If $P_3 = \mathcal{O}$, then the result of the addition is \mathcal{O} .

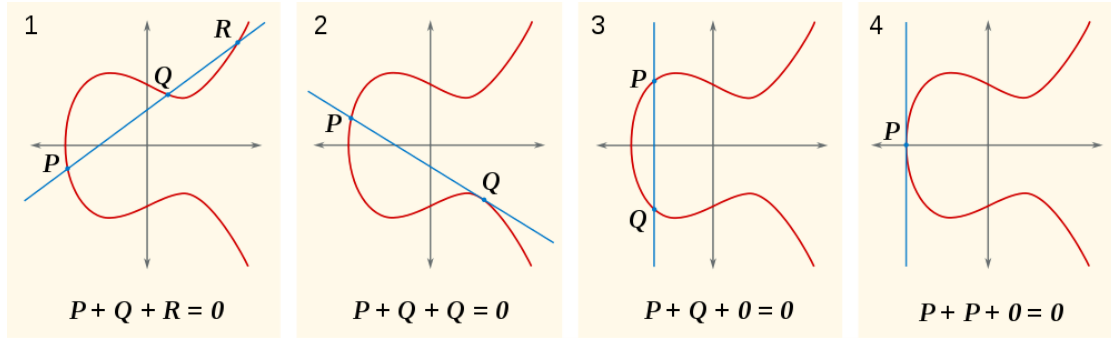


Figure 2.3: Elliptic Curve point addition

(Image by SuperManu, licensed under Creative Commons.)

In addition, we define multiplication of a point $P \in E(\mathbb{Z}_p)$ with a scalar $n \in \mathbb{N}$ as follows:

$$n \cdot P = \underbrace{P + P + \dots + P}_{n \text{ times}} \quad (2.3)$$

That is, the product of a scalar with a point is again a point on our curve.

2.4.2 Groups on Elliptic Curves

For any point $P = (x, y)$, there is unique inverse $-P = (x, -y)$. In addition associativity holds (i.e. for all $P_1, P_2, P_3 \in E(\mathbb{Z}_p)$, $(P_1 + P_2) + P_3 = P_1 + (P_2 + P_3)$). Thus, $E(\mathbb{Z}_q)$ with point addition is a group. [2]

Of particular interest to cryptography are *cyclic* groups on elliptic curves. A (multiplicative) group G is cyclic if there is an element $g \in G$ that generates G , i.e. $G = \langle g \rangle = \{g^n | n \in \mathbb{Z}\}$. Translated to our (additive) groups on elliptic curves, this means that there is a generator point $P \in E(\mathbb{Z}_p)$, such that every point $Q \in E(\mathbb{Z}_p)$ can be written as $Q = nP$ with some $n \in \mathbb{N}$.

In particular, if a group has prime order (i.e. the number of points $\#E(\mathbb{Z}_p)$ on a curve is prime), all points except \mathcal{O} are generators. This follows from the fact that for any point $P \in E(\mathbb{Z}_p)$, its order $\text{ord}(P) = |\langle P \rangle|$ must divide the group order. A prime has exactly two divisors: One (the order of \mathcal{O}) and itself (the order of all other points).

These cyclic elliptic curve groups (or, cyclic subgroups of non-cyclic elliptic curves) are exactly the groups we are interested in for doing actual cryptography. See also [2, p. 321].

2.4.3 Hardness Assumptions

Most ECC schemes are build upon three hardness assumptions: The Discrete Logarithm Problem (DLP), the Decisional Diffie-Hellman Problem (DDHP) and the Computational Diffie-Hellman Problem (CDHP). Given an additive, cyclic group \mathbb{G} with $P \in \mathbb{G}$ a generator, they are stated as follows:

Discrete Logarithm Problem. Given an arbitrary $Q \in \mathbb{G}$, compute an $n \in \mathbb{N}$ such that $nP = Q$.

Computational Diffie-Hellman Problem. Given the triple (P, aP, bP) where $a, b \in \mathbb{N}$ chosen uniformly at random, compute abP .

Decisional Diffie-Hellman Problem. Given two triples (aP, bP, abP) and (aP, bP, Q) where $a, b \in \mathbb{N}$ and $Q \in \mathbb{G}$ chosen uniformly at random, distinguish between the two.

Now, the hardness *assumption* is that, for some groups, these problems are hard to solve, i.e. solving them requires so much time and computational power that it is infeasible. From this, we can build secure asymmetric encryption schemes.

These three problems might all be assumed to be *hard*, but that doesn't mean they are equally so: If, in a certain group \mathbb{G} , the DLP problem is easy, so is CDHP: Just compute a and b , and then use them to calculate abP . And if CDHP is easy w.r.t some \mathbb{G} , so is DDHP: Just compute abP , and compare the third element of each tuple. The inverse is not generally true, i.e. there are groups in which DLP and CDHP are hard to solve, even though DDHP is easy to solve. In that sense, DLP is the hardest and DDHP the easiest of the three. [2] (TODO find out if this is true. [3])

2.4.4 Bilinear Pairings

Introduction

The final primitive needed for pairing-based encryption schemes are bilinear pairings. These are functions mapping two points on (possibly different) elliptic curves to elements of a finite field (*not* another point on a curve).[4]

Let $n \in \mathbb{N}_0$, \mathbb{G}_1 denote an additive abelian group of order n with generator P and identity \mathcal{O} . Let \mathbb{G}_T be another group of order n , this time written multiplicatively with generator g and identity 1. A *bilinear pairing* then is a function $e : \mathbb{G}_1 \times \mathbb{G}_1 \rightarrow \mathbb{G}_T$ with the following properties:

- *Bilinearity.* For all $Q_1, Q_2, Q_3 \in \mathbb{G}_1$, $e(Q_1 + Q_2, Q_3) = e(Q_1, Q_3) + e(Q_2, Q_3)$ and $e(Q_1, Q_2 + Q_3) = e(Q_1, Q_2) + e(Q_1, Q_3)$
- *Non-Degeneracy.* $e(P, P) \neq 1$

Using a bilinear pairing with a group \mathbb{G}_T where the DLP is easy to solve, the DLP in \mathbb{G}_1 can also be solved easily: To find n such that $Q = nP$, compute $e(P, Q) = e(P, nP) = e(P, \underbrace{P + \dots + P}_{n \text{ times}}) = \underbrace{e(P, P) \cdots e(P, P)}_{n \text{ times}} = e(P, P)^n$. Thus, the discrete logarithm of Q with respect to P is the discrete logarithm of $e(P, Q)$ with respect to $g = e(P, P) \in \mathbb{G}_T$. [3]

Bilinear Hardness Assumption

Since the DLP can be easy to solve using pairings, we need to adapt our hardness assumption. Therefore, we define

The *Bilinear Diffie-Hellman Problem* (BDHP): Given P, aP, bP, cP with $a, b, c \in \{1, \dots, n-1\}$ chosen randomly, compute $e(P, P)^{abc}$. [3]

If the BDHP is hard for a pairing e on groups \mathbb{G}_1 and \mathbb{G}_T , this implies that the DLP is hard in both \mathbb{G}_1 and \mathbb{G}_T . [3]

2.5 Different ABE schemes for use in Embedded Devices

3 Challenges encountered during the implementation

3.1 Rust specialties

The borrow checkaaa...

3.2 Lack of Operating System

This means lack of

- allocator and standard library (replaced by `core`, but much less powerful)
 - any dynamically allocated data structures
 - No Vectors and HashMaps Vectors and HashMaps
 - no easy implementation of (access-) trees using recursive enums (as these would become infinitely large when not using indirections)
 - most dependencies depend on `std` in some way, so a lot more work to make all those independent
 - Solution: `heapless` crate and linear representation of access trees (TODO make bounds on size of data types configurable)
- Random Number Generation
 - Problem: need cryptographically secure randomness, but have no OS randomness pool
 - can't just use `/dev/urandom` to get randomness (there is no such thing as a file anyway)
 - standard implementations of `random` resp. `getrandom` crates don't work (rely on `stdlib`)
 - need own randomness source
 - `nrf50` series provides hardware RNG, but we need Rust to be able to interface with this

- probably need own implementation of `getrandom` crate for Zephyr OS or bare-metal nrf50
- Unit testing abilities

3.3 Performance Limitations

That is, CPU speed (64MHz) and RAM size (..KiB?).

- probably too slow for computing actual pairings
- Solution 1: Scheme without pairings (Yao et al 2015)
- Solution 2: Still do pairings, but hyper-optimize (see TinyPBC on AVR)

List of Figures

1.1	Example drawing	1
1.2	Example plot	2
1.3	Example listing	2
2.1	Keys used for encryption and decryption in different classes of encryption schemes	4
2.2	Example for a $(5, 4)$ -threshold scheme with $s = 12$ and $p(x) = 12 + 7x - 6x^2 + x^3$. As $p(x)$ has degree three, at least four of the five shares are required to reconstruct s	6
2.3	Elliptic Curve point addition (Image by SuperManu, licensed under Creative Commons.)	8

List of Tables

1.1	Example table	1
-----	-------------------------	---

Bibliography

- [1] A. Shamir, "How to share a secret," *Communications of the ACM*, vol. 22, no. 11, pp. 612–613, Nov. 1979, ISSN: 0001-0782, 1557-7317. DOI: 10.1145/359168.359176.
- [2] J. Katz and Y. Lindell, *Introduction to modern cryptography*, second edition, ser. Chapman Hall, CRC cryptography and network security. Boca Raton ; London ; New York: CRC Press, 2015, xx, 583 SeitenIllustrationen, ISBN: 978-1-4665-7026-9 1-4665-7026-1.
- [3] A. Menezes, "An introduction to pairing-based cryptography," in *Contemporary Mathematics*, I. Luengo, Ed., vol. 477, Providence, Rhode Island: American Mathematical Society, 2009, pp. 47–65, ISBN: 978-0-8218-3984-3 978-0-8218-8156-9. DOI: 10.1090/conm/477/09303.
- [4] I. Blake, G. Seroussi, and N. Smart, Eds., *Advances in elliptic curve cryptography*, Lecture note series. Cambridge [u.a.]: Cambridge University Press, 2005, XVI, 281 S. ISBN: 0-521-60415-X.